

# Kathará

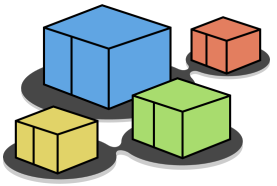
## data centers' architecture

<b>Version</b>	3.1
<b>Author(s)</b>	L. Ariemma, G. Di Battista, M. Patrignani, M. Scazzariello, T. Caiazzì
<b>E-mail</b>	contact@kathara.org
<b>Web</b>	<a href="http://www.kathara.org/">http://www.kathara.org/</a>
<b>Description</b>	Data Centers' Routing: Multipath, Fat-Trees



# copyright notice

- all the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as “material”) are protected by copyright
- this material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide
- this material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes
- any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement
- this copyright notice must always be redistributed together with the material, or its portions



# Multi-Path

in a nutshell



# Multi-Path

- data center architectures allow to establish several paths between pairs of nodes
- Multi-Path routing is used to
  - distribute and balance the traffic among different paths
  - increase the fault tolerance



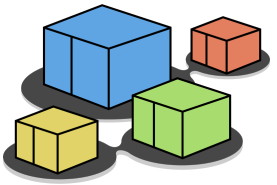
# Multi-Path

- allows to have more than one next-hop for the same prefix in the FIB (Forwarding Information Base)
- needs kernel support
  - the main OSes and routers support it
- different usage policies can be applied
- packets of the same *flow* should use the same path
  - reordering packets is computationally heavy



# Equal-Cost Multi-Path (ECMP)

- a specific approach to Multi-Path
- exploits paths to the same destination that have the same “cost”
  - e.g., same IGP cost, same number of hops
- allows to choose among the available paths in a uniform way
- often used in Fat-Tree data centers
  - we’ll get there in a couple of slides



# how packets are forwarded

- for each packet, the kernel decides the FIB entry to be used
- different policies are allowed
  - hash-based policy
    - Layer-3 hash (src IP, dst IP)
    - Layer-4 hash (src IP, dst IP, src port, dst port, protocol)
  - round-robin



# Hyper-scale Data Centers

main concepts






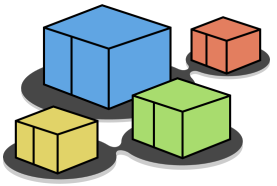
# a high-level architecture

- Internet connections
  - multiple redundant high-speed fiber optic links
- fabric
  - infrastructure designed to transport packets between servers and between servers and Internet
- server farm
  - host applications and services



# a high-level architecture





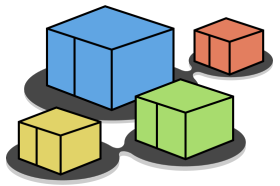
# fabric architecture

why Fat-Trees?



# overview – fabric architecture

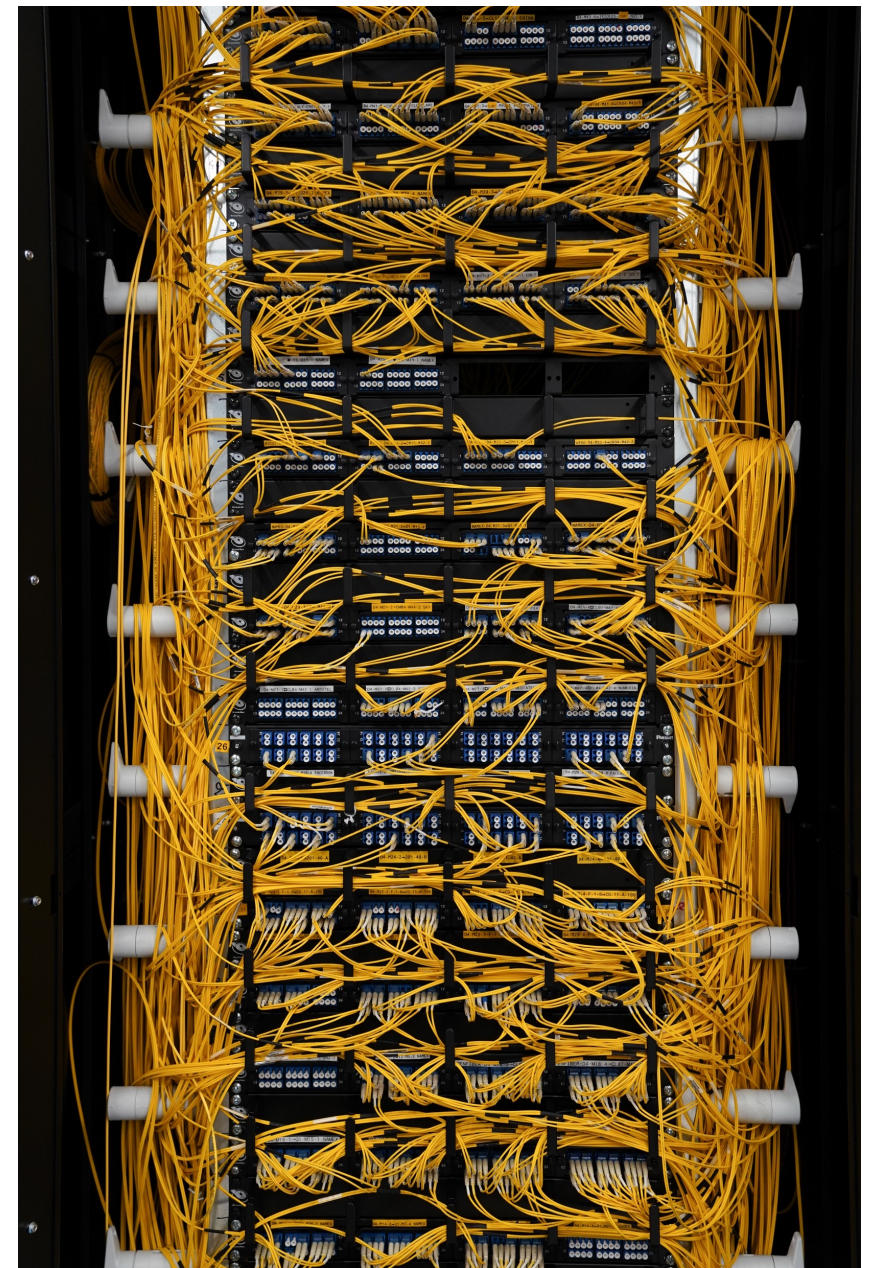
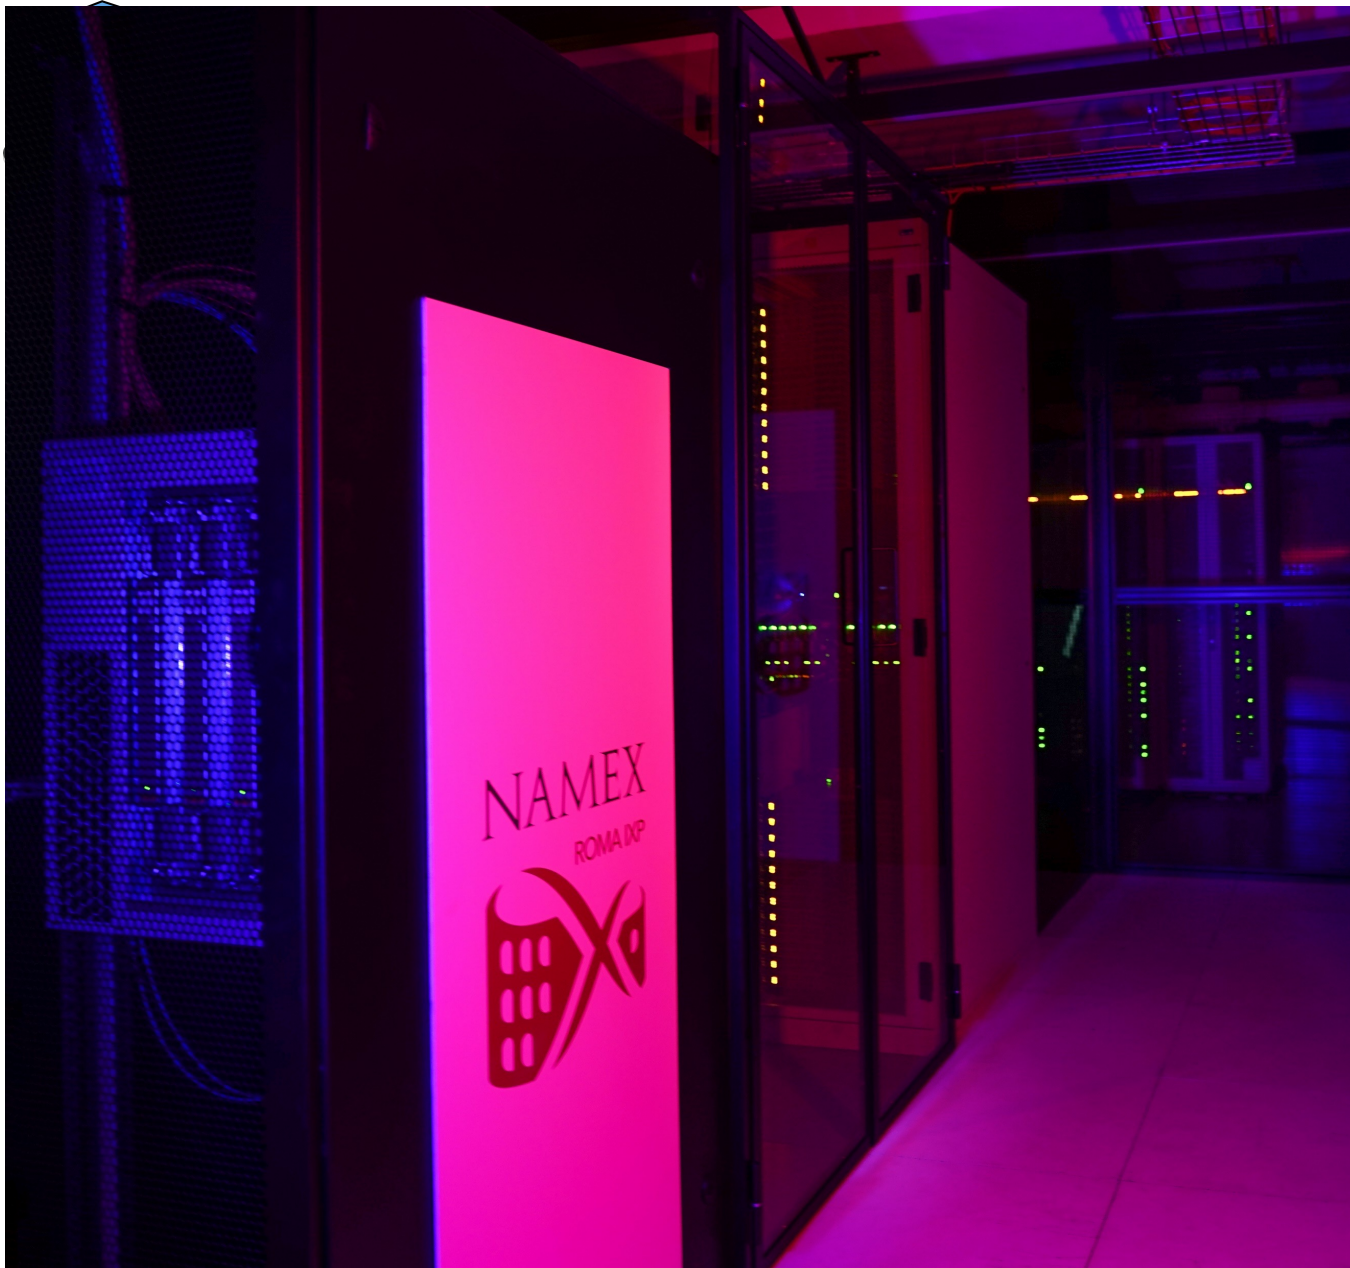
- real-life data centers
- service model
- data center network
  - traffic flow directions
  - requirements
  - components
  - topologies



# real data center information – figures

- links are fiber optic links
- each link in the fabric is at least at 100Gb/s
- each switch is at least of 64 ports
- about 5,000 switches/routers
- about 60,000 servers







# data center failures – yearly report

statistics by Google (2008) in a *portion* of data center composed of 1,800 servers:

- 1,000 individual machine (switch/router or server) failures
- thousands of hard drive failure
- 1 power distribution unit failure
  - down for 500/1000 machines for 6 hours
- 1 cluster complete rewire (not simultaneously)





# data center service models

## ■ three different service models

### ■ on-premise data centers

- built, owned and operated by a company
- often housed in a building of the organization

### ■ colocation data centers

- built and owned by a company that rents space within the data center to other companies
  - the hosting company manages the infrastructure (building, cooling, bandwidth, security etc.)
  - the hosted companies provides the components, including servers, storage, and firewalls

### ■ cloud data centers

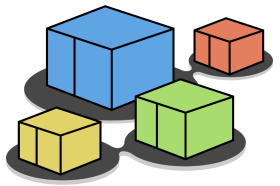
- applications and data are hosted by a cloud service provider such as Amazon Web Services (AWS), Microsoft Azure, and Google Cloud Platform (GCP)



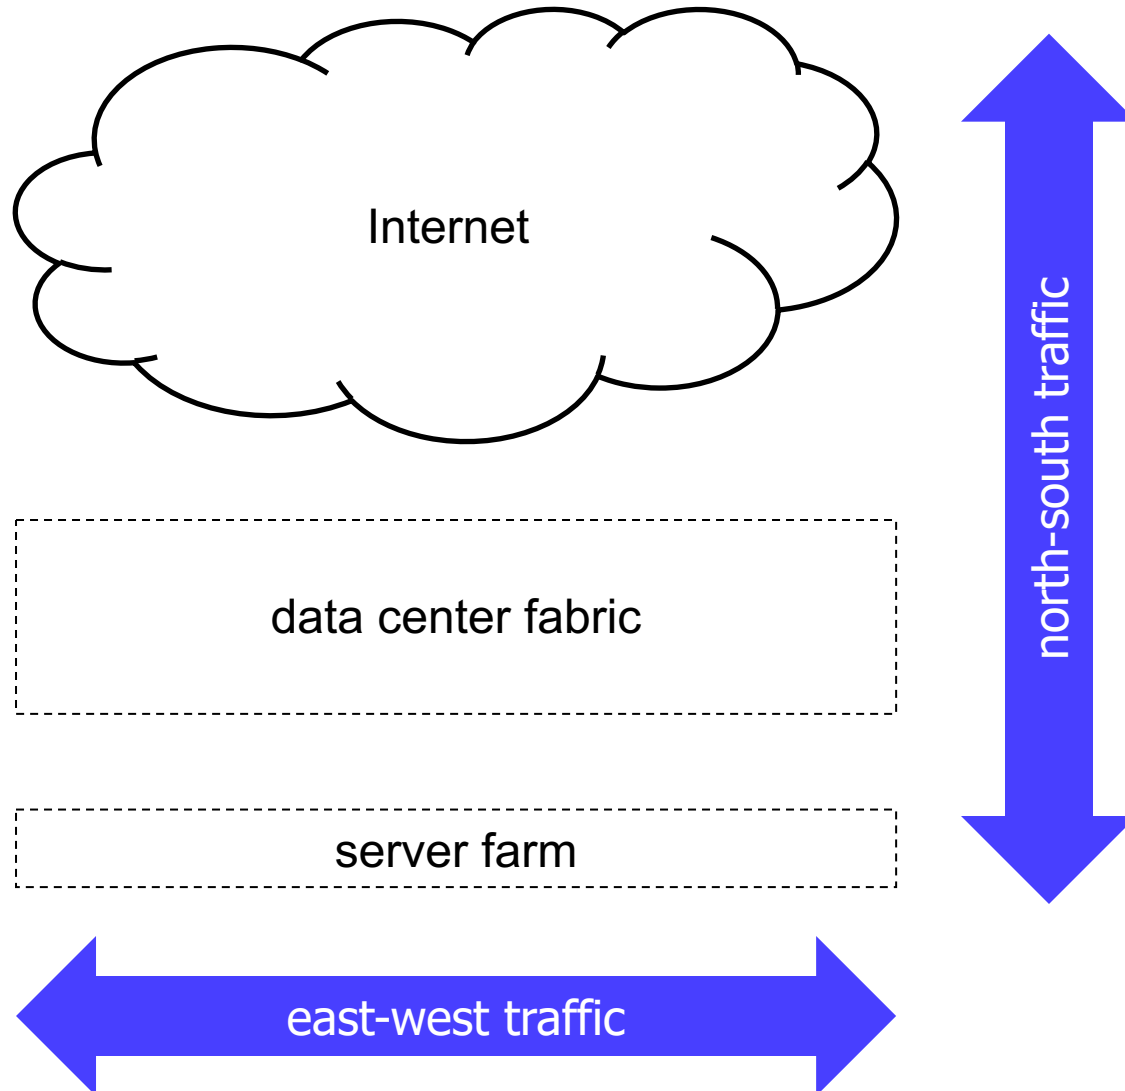


# traffic flows inside a data center

- two types of flows can be identified
  - traffic exiting or entering the data center
    - also called “north-south traffic”
    - data sent/received via Internet
  - server-to-server communications
    - also called “east-west traffic”: in data center schemata servers are typically drawn side-by-side
    - primarily, supports micro-services and distributed architectures



# traffic flows inside a data center





# data center network requirements

- support high-bandwidth server-to-server communication
  - applications that rely on cluster computations, such as Hadoop or Spark, can involve hundreds or thousands of servers
  - customer's containers/virtual machines (VMs) are distributed across multiple servers but need to communicate seamlessly
  - microservice architectures heavily rely on server-to-server communication
- scale
  - data centers range from a few hundred to a hundred thousand servers in a single physical location
- resilience
  - data center applications are designed to work in presence of failures



# data center network components

- data center nodes can be connected by using two network component types
  - specialized hardware
    - proprietary hardware that can scale to clusters of thousands of nodes with high bandwidth
      - e.g., Google Jupiter and InfiniBand switches
    - expensive option
  - commodity switches and routers
    - cheap option
    - widely adopted, this is the option we consider in the following



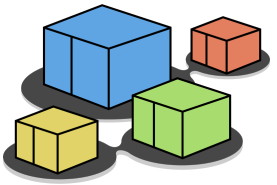
# topology requirements

- scalability
  - it should be possible to expand the data center
  - in other words, it must be possible to buy and deploy hardware similar to the one already deployed
- bandwidth
  - hosts in the fabric should communicate with each other using the full bandwidth of their NICs



# data center network topologies

- several topologies have been proposed
  - Clos
  - Fat-Tree
  - VL2
  - Jellyfish
  - Xpander
  - DCell
  - ...



# Clos topology

- invented by Edson Erwin in 1938 to address scalability issues in telephone networking
  - formalized by Charles Clos in 1953
- the original problem was allowing  $N$  contemporary connections from  $N$  input lines to  $N$  output lines without using a single crossbar switch



# Fat-Tree topology

- Fat-Trees were originally introduced by Charles Leiserson in 1985
- the Fat-Tree is a special case of a Clos
- typical Fat-Tree architectures today consist of three level of nodes (switches/routers)
- high redundancy
- constant *bisection* of the available bandwidth
- typically, nodes have the same characteristics
  - easier to stock spare equipments






# Fat-Tree parameters

- the Fat-Tree is a modular topology
- two parameters define a Fat-Tree
  - *radix* of the nodes (even number, denoted by  $2K$ )
    - number of available ports
  - *redundancy factor* (denoted by  $R$ )
- usually, if the radix of a node is  $2K$  it has  $K$  connections towards the north and  $K$  connections towards the south
  - different choices are possible, but not considered here

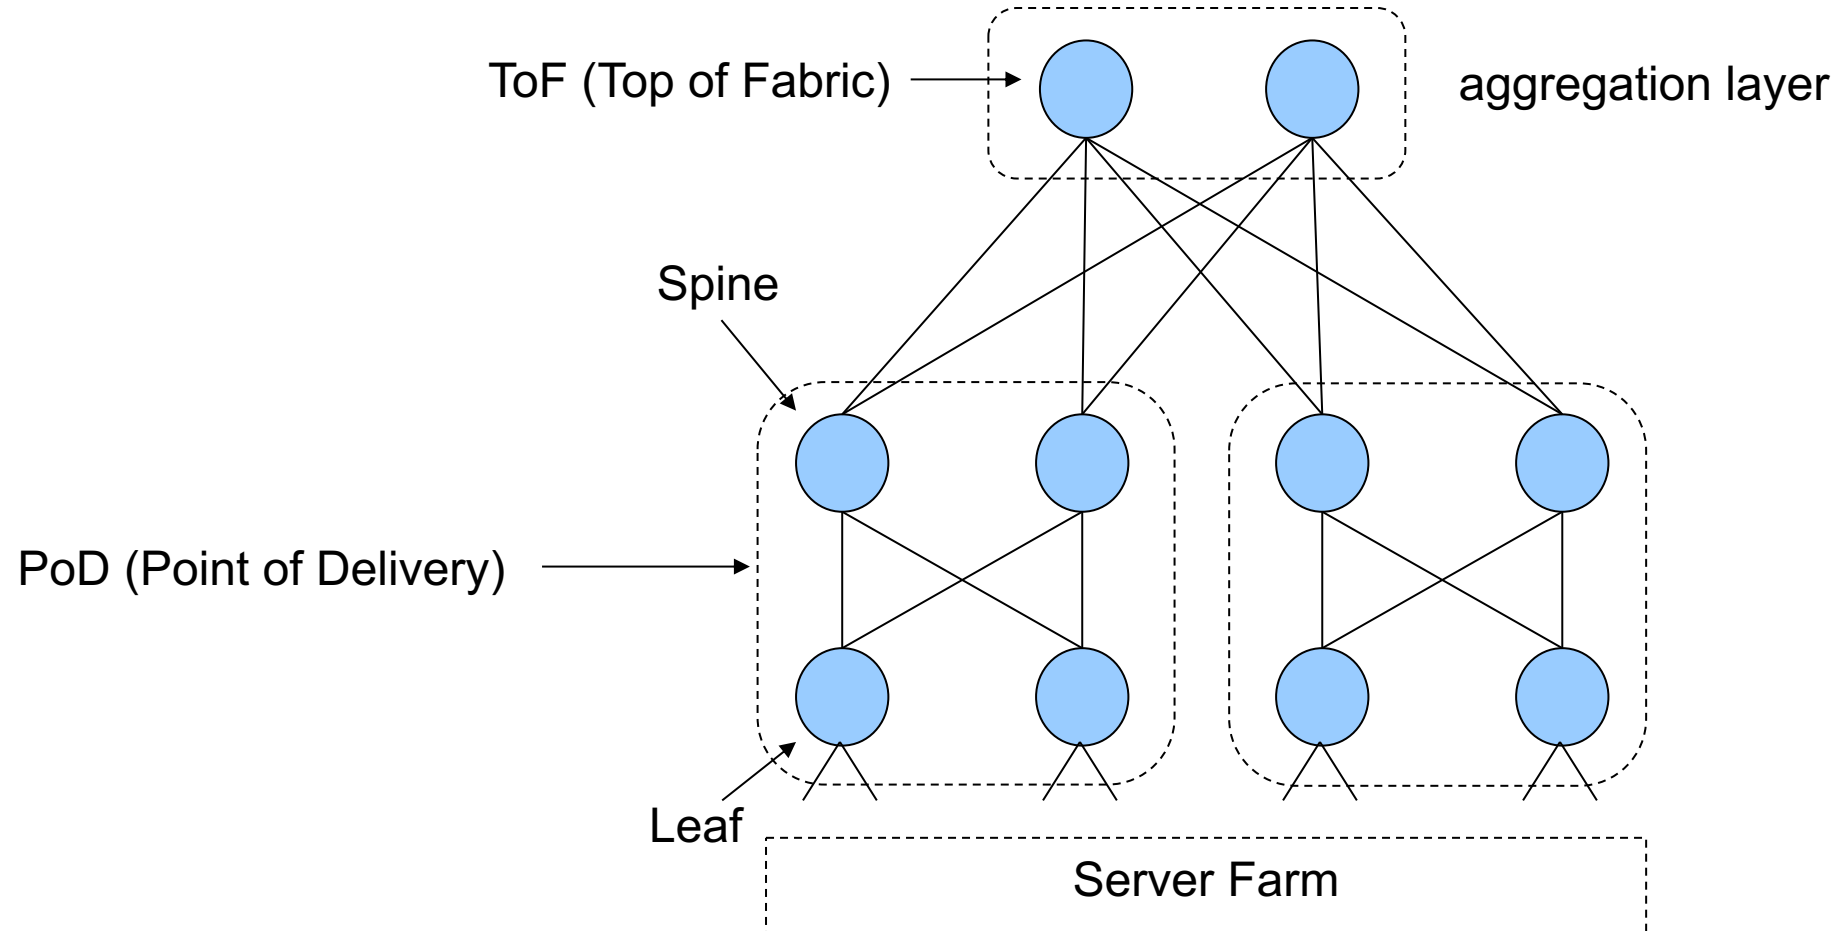


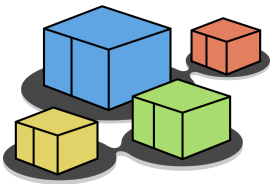
# Fat-Tree nodes

- Leaf
  - node connected to the server farm
- Spine
  - node north of Leaves and south of ToF nodes
- Point of Delivery (PoD)
  - set of fully interconnected Leaves and Spines
- Top of Fabric (ToF)
  - set of top nodes that provide inter-PoD communication





# an example of Fat-Tree: $FT(K=2, R=2)$

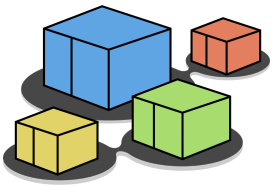




# redundancy factor (R)


- number of links between a ToF node and a PoD
- allow to connect more PoDs reducing the redundancy





# multi-plane Fat-Tree

- when  $K \neq R$  the Fat-Tree is called *multi-plane*
- the ToF nodes are partitioned in sets called *planes*





# example of a multi-plane Fat-Tree


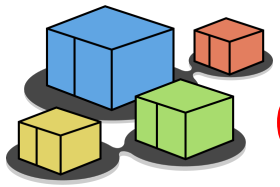

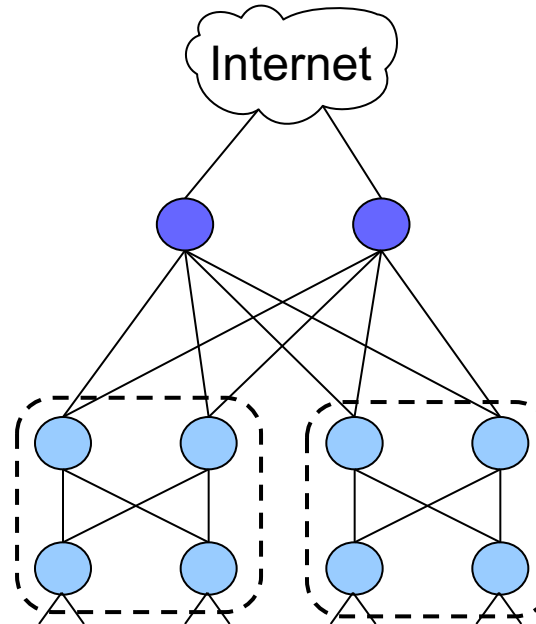


diagram from the  
IETF RIFT Draft



# connecting a Fat-Tree to the Internet

- two strategies
  - usage of a dedicated PoD
  - usage of ToFs





# bibliography and further readings

- [Caiazzì '22] Caiazzì, Scazzariello, Alberro, Ariemma, Castro, Grampin, Di Battista, "Sibyl: a Framework for Evaluating the Implementation of Routing Protocols in Fat-Trees", NOMS 2022
- [Caiazzì '21] Caiazzì, Scazzariello, Ariemma, "VFTGen: a Tool to Perform Experiments in Virtual Fat Tree Topologies", IM 2021
- [Caiazzì '19] Caiazzì, "Software Defined Data Centers: methods and tools for routing protocol verification and comparison", Ms. Thesis, Roma Tre University, 2019
- [Dutt '17] Dutt, "BGP in the Data Center", O'Reilly, 2017
- [RFC-7938] Lapukhov, Premji, "Use of BGP for Routing in Large-Scale Data Centers" Internet Engineering Task Force (IETF) Request for Comments: 7938