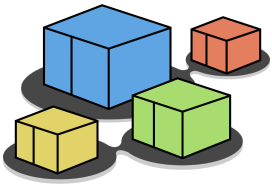


Kathará

kathara lab

rip with FRRouting

Version	1.1
Author(s)	G. Di Battista, M. Patrignani, M. Pizzonia, L. Ariemma, M. Scazzariello, T. Caiazzi
E-mail	contact@kathara.org
Web	http://www.kathara.org/
Description	experiences with the ripv2 distance vector routing protocol – derives from kathara rip lab ver. 1.2 which, in turns, derives from netkit rip lab ver. 2.4




copyright notice

- All the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as “material”) are protected by copyright.
- This material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide.
- This material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes.
- Information contained in this material cannot be used within network design projects or other products of any kind.
- Any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement.
- The authors assume no responsibility about this material and provide this material “as is”, with no implicit or explicit warranty about the correctness and completeness of its contents, which may be subject to changes.
- This copyright notice must always be redistributed together with the material, or its portions.



the RIP protocol

- RIP is a routing protocol
- RIPS is a distance-vector routing protocol
 - approach: send all your information to a few
 - update your routing information based on what you hear
- in this lab we will see an example of RIPv2 protocol on the FRRouting Suite



sample `frr.conf` configuration file

```
root@pc1:~$ cat /etc/frr/frr.conf
```

```
!  
! FRRouting configuration file
```

talk rip on some interface

```
!  
!  
! RIP CONFIGURATION
```

```
!  
router rip  
redistribute connected  
network 100.1.0.0/24
```

redistribute to rip neighbors'
information about all directly
connected subnets


```
!  
log file /var/log/frr/frr.log
```

send rip multicast
packets to interfaces
falling into this
prefix





about `redistribute connected`

- by default (i.e., without further configuration) RIP already propagates information about all directly connected subnets attached to RIP-speaking interfaces
- `redistribute connected` forces RIP to propagate information about all connected subnets
- the semantic of `redistribute connected` applies to all routing protocols
- the default behavior does not
 - some protocols (e.g., bgp) are lazier, and do not propagate anything unless explicitly told to do so




small network connected to Internet







the involved ip subnets

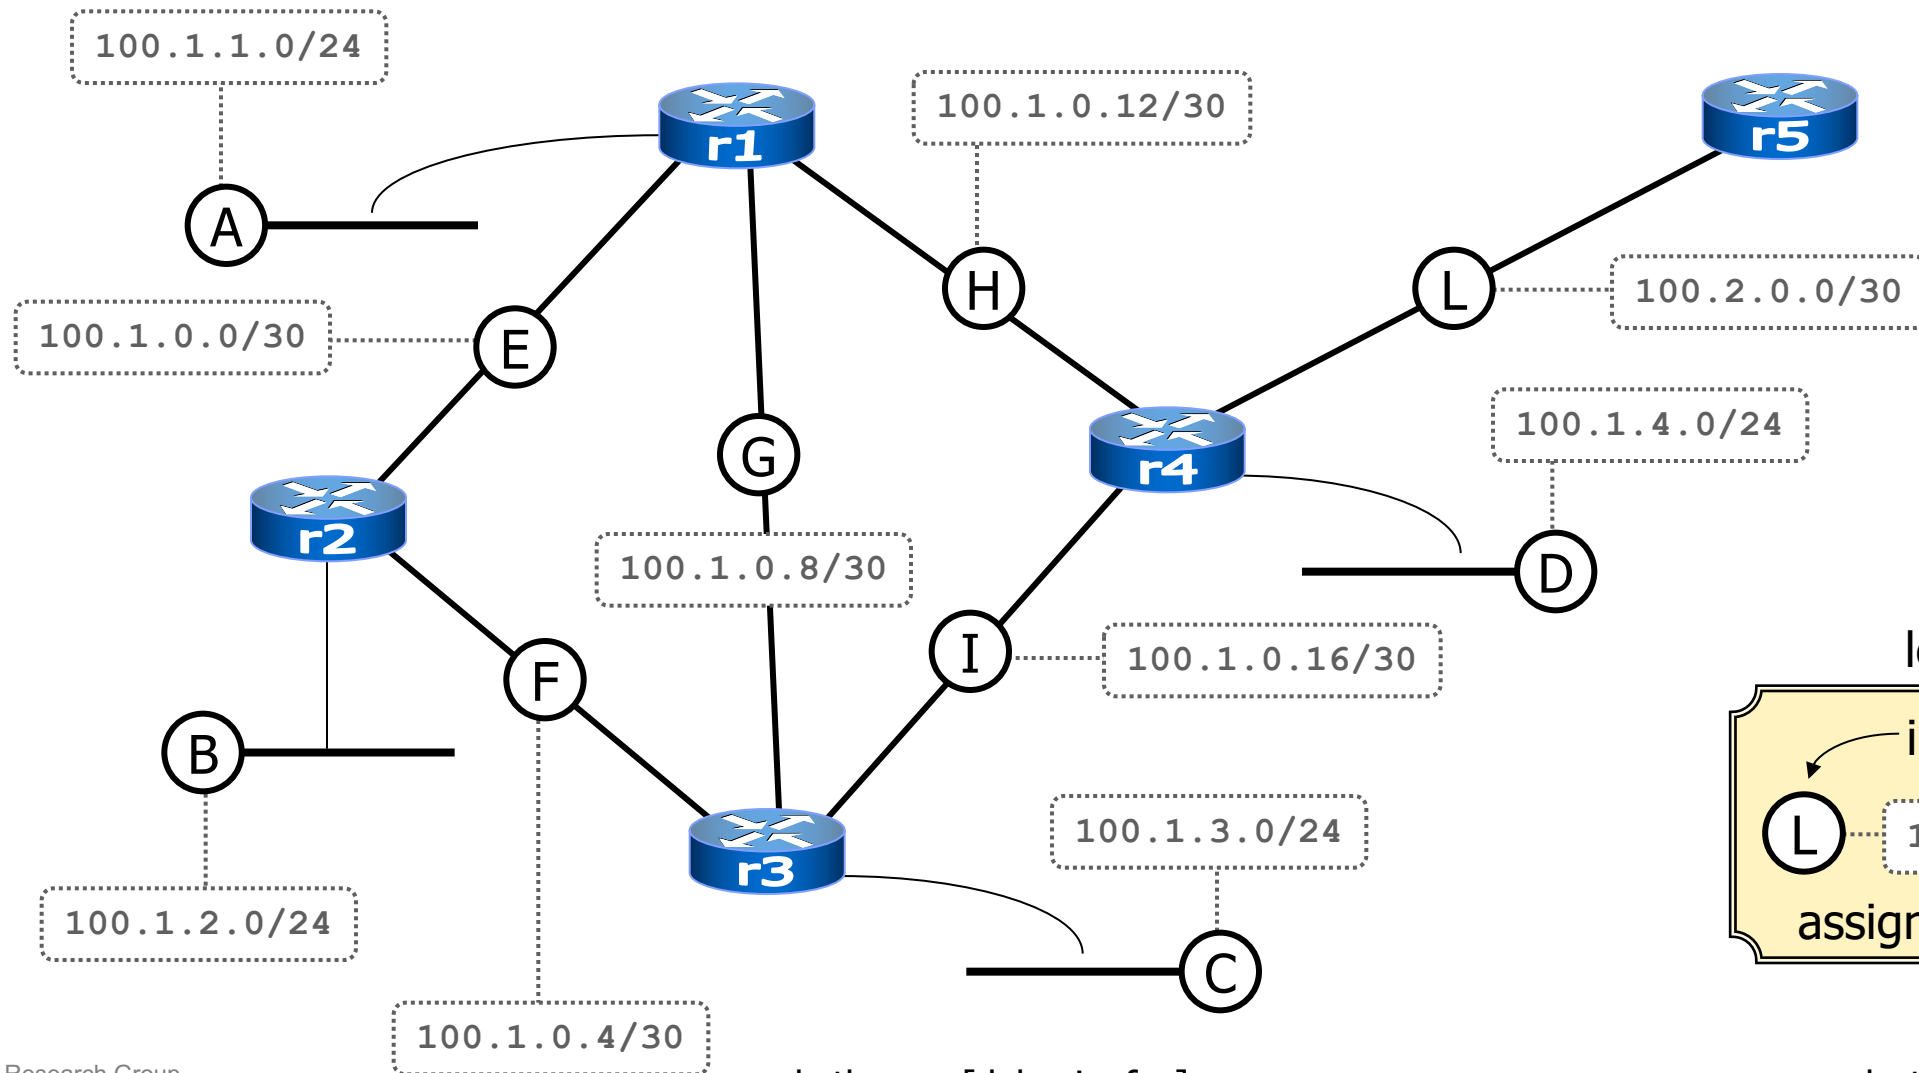



legend

each circle  represents a subnet

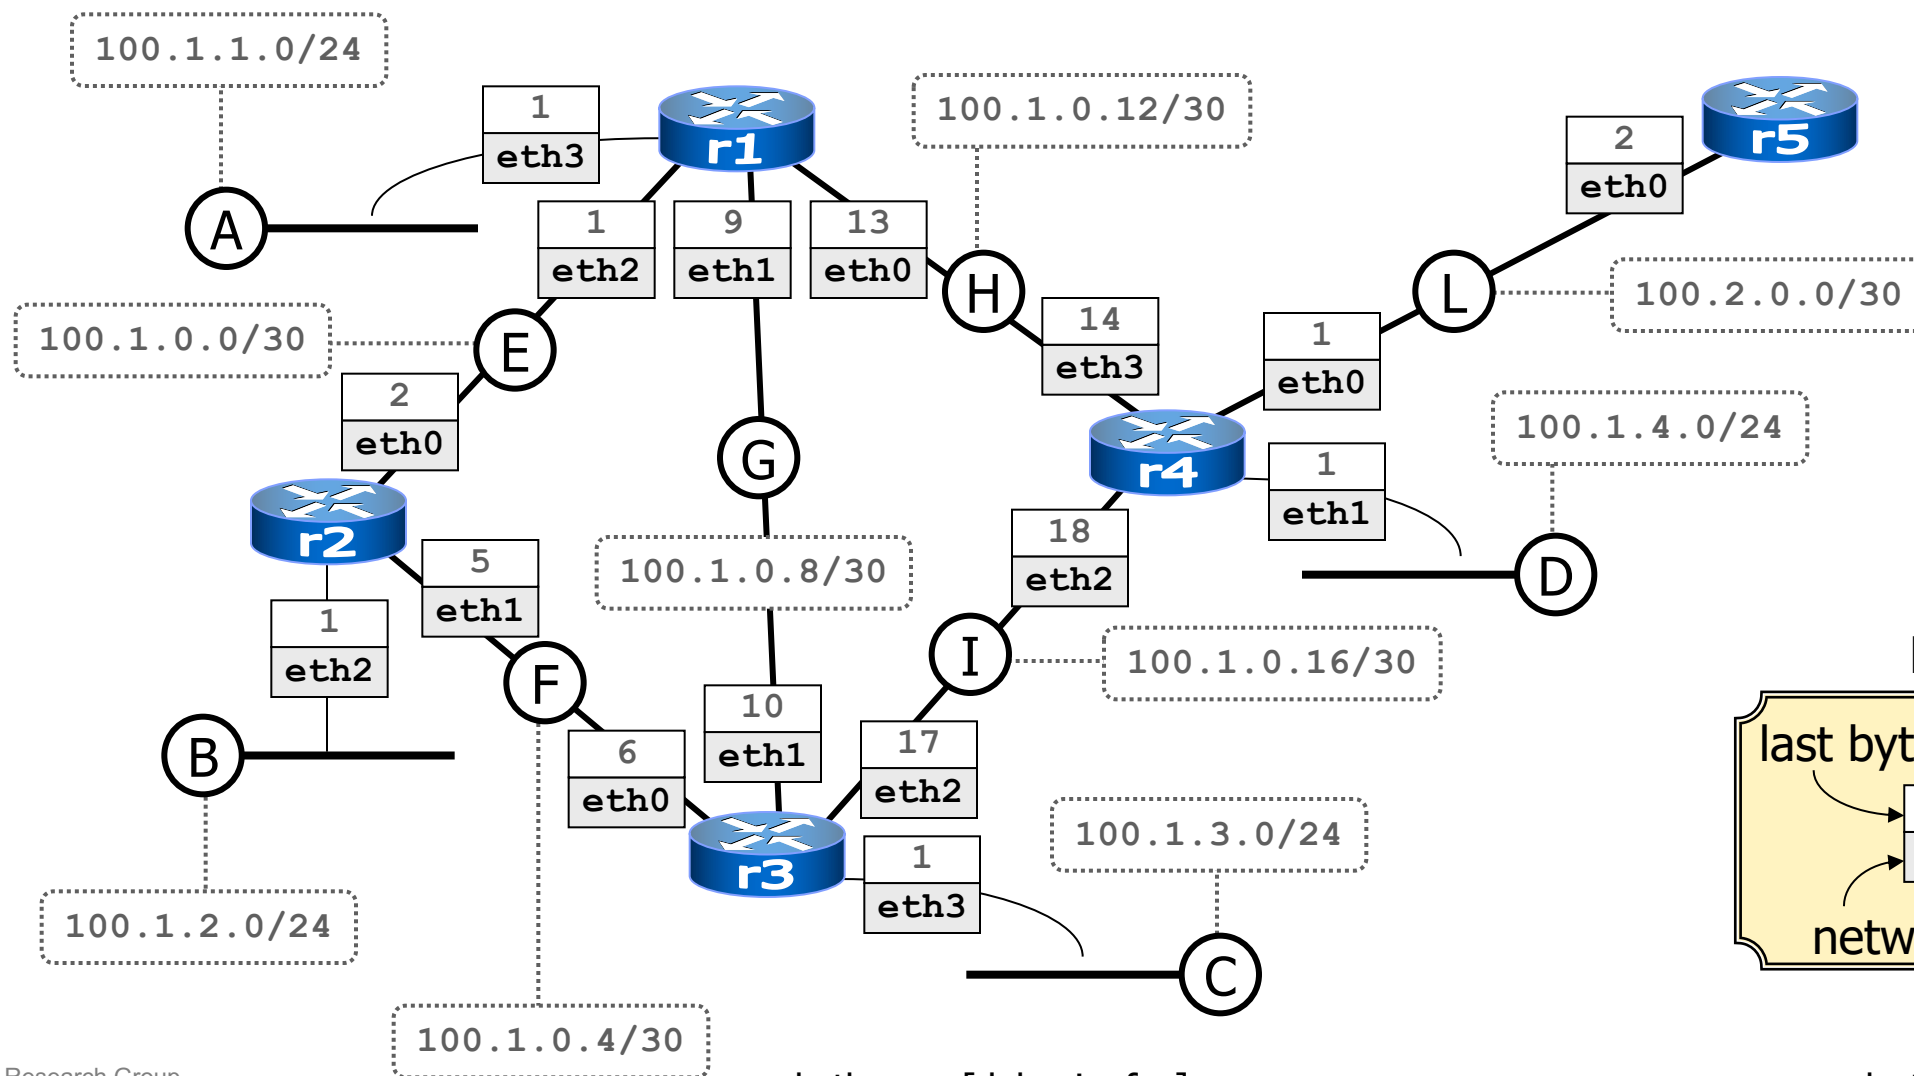


assigning ip numbers to subnets






assigning ip numbers to interfaces



legend





launching the lab

```
root@localhost:~$ cd kathara-lab_rip  
root@localhost:~/kathara-lab_rip$ kathara 1start
```

- the lab configuration is such that the frr routing daemon is not automatically started

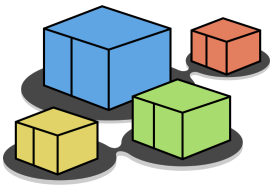


checking connectivity

- towards a directly connected destination

```
r4:~# ping 100.1.0.13
PING 100.1.0.13 (100.1.0.13) 56(84) bytes of data.
64 bytes from 100.1.0.13: icmp_seq=1 ttl=64 time=1.23 ms
64 bytes from 100.1.0.13: icmp_seq=2 ttl=64 time=0.592 ms
64 bytes from 100.1.0.13: icmp_seq=3 ttl=64 time=0.393 ms

--- 100.1.0.13 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time
2032ms
rtt min/avg/max/mdev = 0.393/0.741/1.238/0.360 ms
r4:~#
```

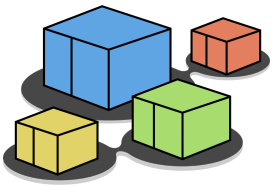


checking connectivity

- towards a remote destination

```
r4:~# ping 100.1.2.1
connect: Network is unreachable
r4:~#
```

- what's going on?



examining the kernel routing table

```
r4
root@r4:/# route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use
Iface
100.1.0.12       0.0.0.0          255.255.255.252 U        0      0      0 eth3
100.1.0.16       0.0.0.0          255.255.255.252 U        0      0      0 eth2
100.1.4.0        0.0.0.0          255.255.255.0   U        0      0      0 eth1
100.2.0.0        0.0.0.0          255.255.255.252 U        0      0      0 eth0
root@r4:/#
```

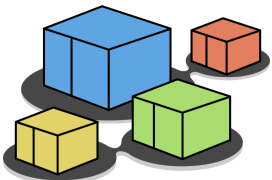
- since no routing daemon is currently running, only directly connected destinations are known to the router



starting the routing daemons

- on each router (but r5) issue the following command:

```
r4:~# /etc/init.d/frr start  
[ ok ] Started watchfrr.  
r4:~#
```



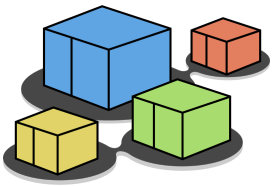
checking connectivity (again)

- towards a remote destination

```
r4:~# ping 100.1.2.1
PING 100.1.2.1 (100.1.2.1) 56(84) bytes of data.
64 bytes from 100.1.2.1: icmp_seq=1 ttl=63 time=0.743 ms
64 bytes from 100.1.2.1: icmp_seq=2 ttl=63 time=0.875 ms
64 bytes from 100.1.2.1: icmp_seq=3 ttl=63 time=0.685 ms

--- 100.1.2.1 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time
2005ms
rtt min/avg/max/mdev = 0.685/0.767/0.875/0.085 ms
r4:~#
```

- after a while, all remote destinations are reachable



checking the routing table

- the routing table is now updated

```
r4
root@r4:/# route
Kernel IP routing table
Destination        Gateway            Genmask           Flags Metric Ref    Use
Iface
100.1.0.0          100.1.0.13        255.255.255.252  UG        20     0      0
eth3
100.1.0.4          100.1.0.17        255.255.255.252  UG        20     0      0
eth2
100.1.0.8          100.1.0.17        255.255.255.252  UG        20     0      0
eth2
100.1.0.12         0.0.0.0           255.255.255.252  U         0      0      0
eth3
100.1.0.16         0.0.0.0           255.255.255.252  U         0      0      0
eth2
100.1.1.0          100.1.0.13        255.255.255.0    UG        20     0      0
eth3
100.1.2.0          100.1.0.17        255.255.255.0    UG        20     0      0
eth2
```




a look at ripv2 packets

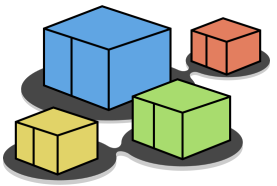
■ let's sniff rinv2 packets

```
r4:~# tcpdump -i eth2 -v -n -s 1518
```

display packet details
(enable full protocol decoding)

don't resolve numbers
to names

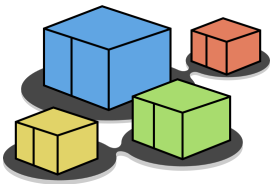
sniff entire ethernet
packets (by default, only
the first 68 bytes are
captured)



a look at ripv2 packets

■ let's sniff ripv2 packets

```
r4
r4:~# tcpdump -i eth2 -v -n -s 1518
tcpdump: listening on eth2, link-type EN10MB (Ethernet), capture size 1518
bytes
16:47:48.333986 IP (tos 0x0, ttl 1, id 0, offset 0, flags [DF], length:
152) 100.1.0.17.520 > 224.0.0.9.520: [udp sum ok]
    RIPv2, Response, length: 124, routes: 6
        AFI: IPv4:      100.1.0.0/30, tag 0x0000, metric: 2, next-hop:
self
        AFI: IPv4:      100.1.0.4/30, tag 0x0000, metric: 1, next-hop:
self
        AFI: IPv4:      100.1.0.8/30, tag 0x0000, metric: 1, next-hop:
self
        AFI: IPv4:      100.1.1.0/24, tag 0x0000, metric: 2, next-hop:
self
        AFI: IPv4:      100.1.2.0/24, tag 0x0000, metric: 2, next-hop:
self
        AFI: IPv4:      100.1.3.0/24, tag 0x0000, metric: 1, next-hop:
self
```



a traceroute

▼ r4



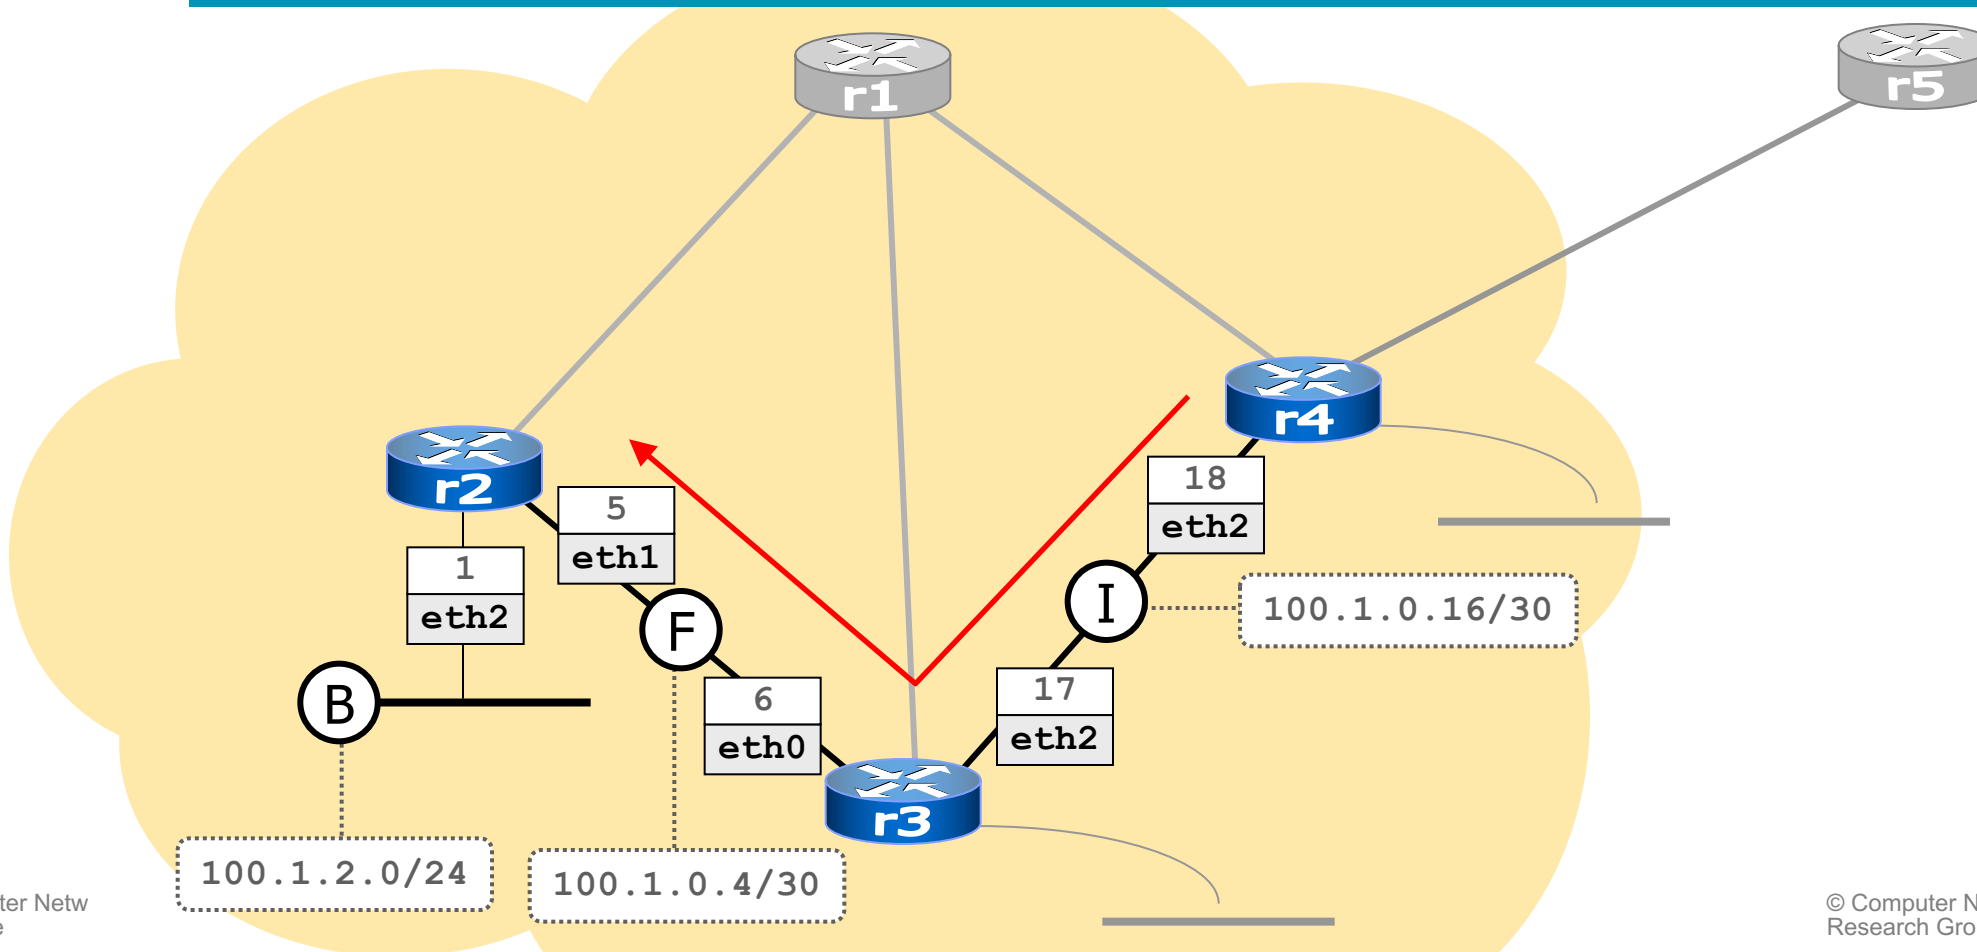
```
r4:~# traceroute 100.1.2.1
```

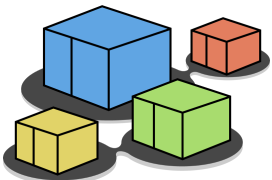
```
traceroute to 100.1.2.1 (100.1.2.1), 64 hops max, 40 byte packets
```

```
1 100.1.0.17 (100.1.0.17) 10 ms 3 ms 1 ms
```

```
2 100.1.2.1 (100.1.2.1) 15 ms 1 ms 1 ms
```

```
r4:~# █
```






inspecting the rip routing table

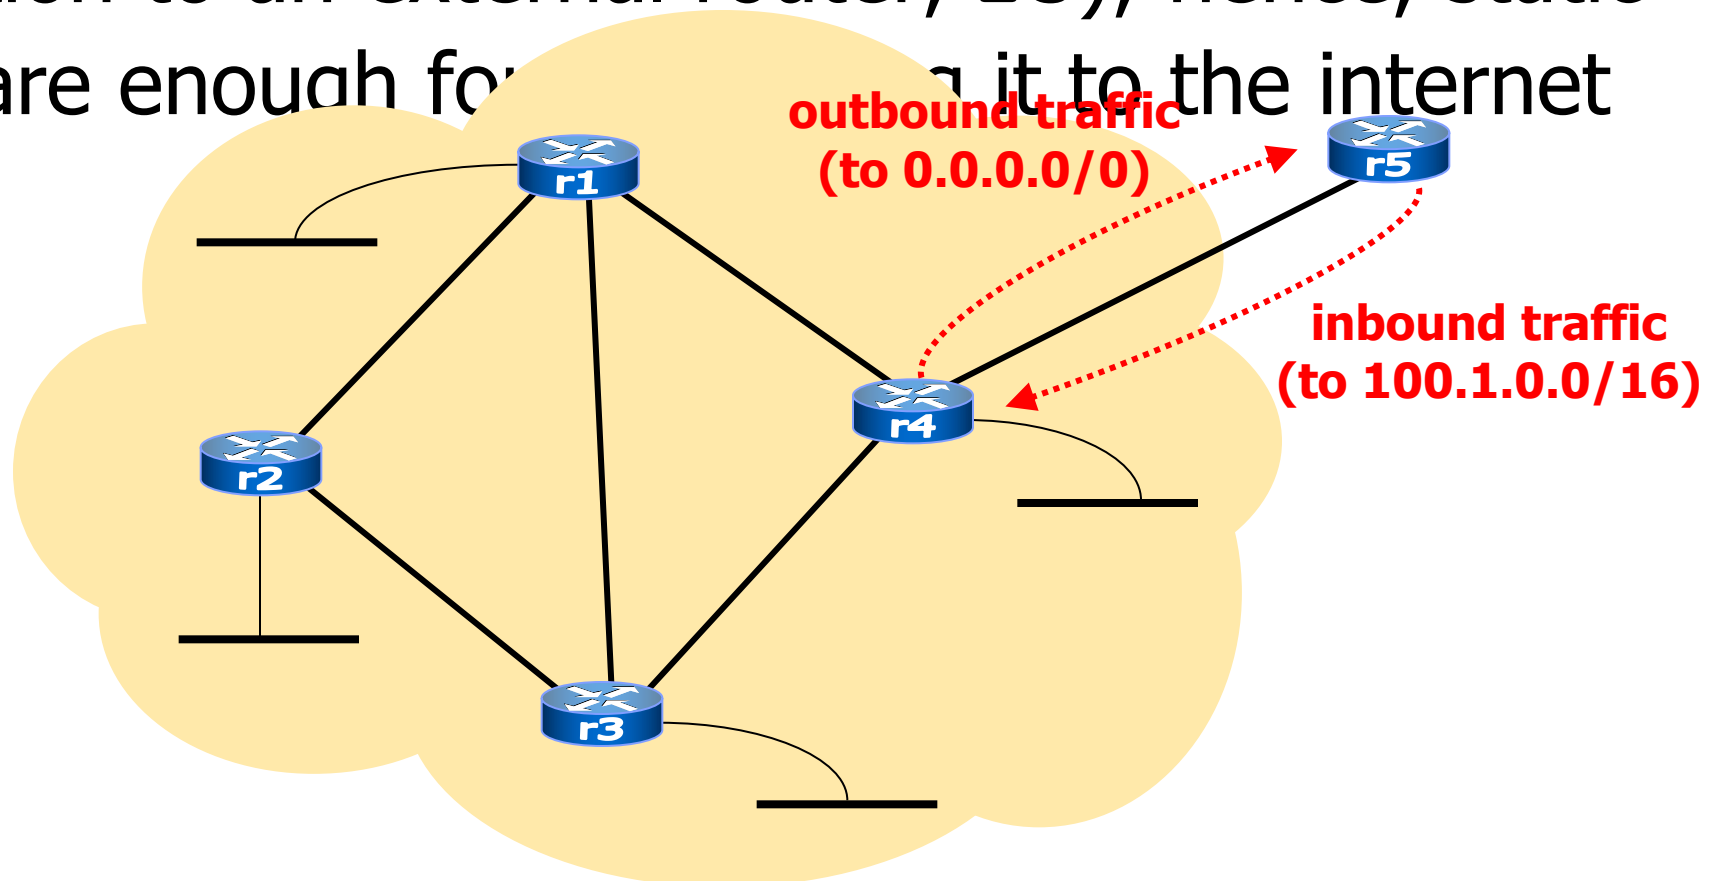
```
▼ r4
root@r4:/# vtysh
r4-frr# show ip rip
Codes: R - RIP, C - connected, S - Static, O - OSPF, B - BGP
Sub-codes:
      (n) - normal, (s) - static, (d) - default, (r) - redistribute,
      (i) - interface

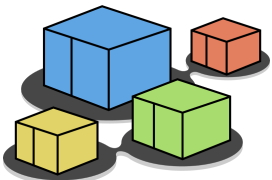
      Network                Next Hop                Metric From                Tag Time
R(n) 100.1.0.0/30            100.1.0.13              2 100.1.0.13              0
02:47
R(n) 100.1.0.4/30            100.1.0.17              2 100.1.0.17              0
02:37
R(n) 100.1.0.8/30            100.1.0.17              2 100.1.0.17              0
02:37
C(i) 100.1.0.12/30           0.0.0.0                 1 self                    0
C(i) 100.1.0.16/30           0.0.0.0                 1 self                    0
R(n) 100.1.1.0/24            100.1.0.13              2 100.1.0.13              0
02:47
R(n) 100.1.2.0/24            100.1.0.17              3 100.1.0.17              0
02:37
R(n) 100.1.3.0/24            100.1.0.17              2 100.1.0.17              0
02:37
```



static routing

- our network is a **stub network** (i.e., it has just one connection to an external router, r5); hence, static routes are enough for connecting it to the internet






adding a static route to r5

```
r5:~# route add -net 100.1.0.0/16 gw 100.2.0.1
r5:~# ping 100.1.2.1
PING 100.1.2.1 (100.1.2.1) 56(84) bytes of data.
64 bytes from 100.1.2.1: icmp_seq=1 ttl=62 time=24.1 ms
64 bytes from 100.1.2.1: icmp_seq=2 ttl=62 time=1.11 ms

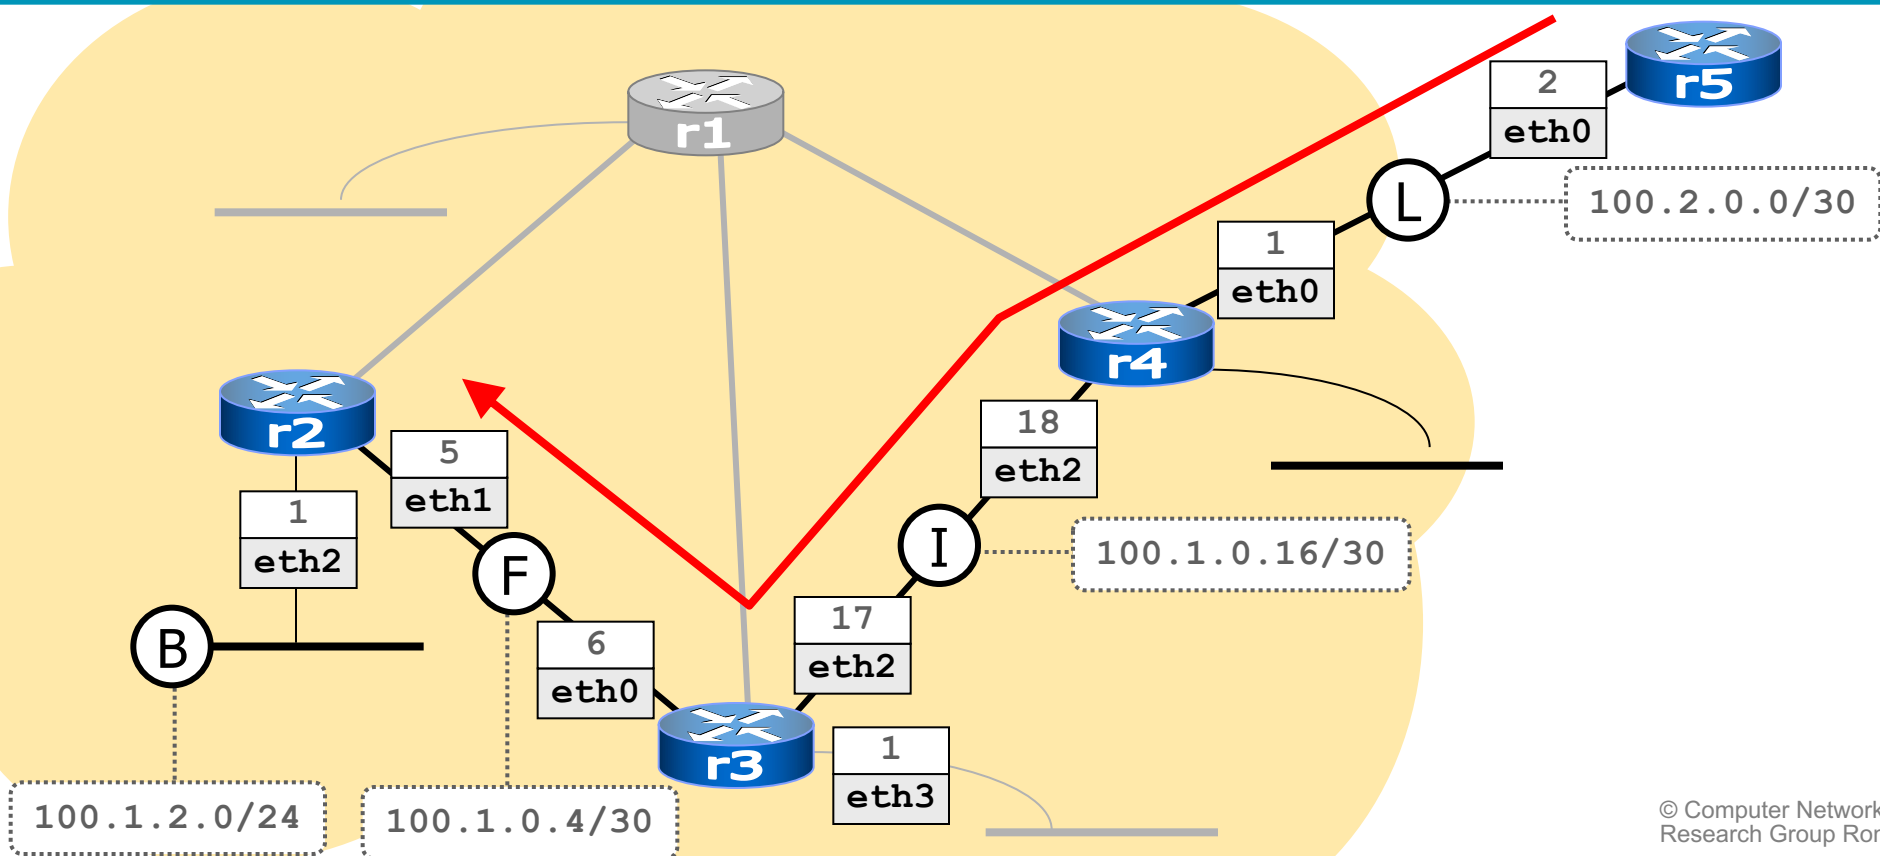
--- 100.1.2.1 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1023ms
rtt min/avg/max/mdev = 1.117/12.634/24.151/11.517 ms
r5:~#
```




checking connectivity

▼ r5

```
r5:~# traceroute 100.1.2.1
traceroute to 100.1.2.1 (100.1.2.1), 64 hops max, 40 byte packets
 1 100.2.0.1 (100.2.0.1) 75 ms 1 ms 2 ms
 2 100.1.0.17 (100.1.0.17) 7 ms 1 ms 1 ms
 3 100.1.2.1 (100.1.2.1) 24 ms 3 ms 1 ms
r5:~# █
```






configuring r4

■ step 1: configuring the default route

```
r4:~# route add default gw 100.2.0.2
root@r4:/# route
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use
default	100.2.0.2	0.0.0.0	UG	0	0	0 eth0
100.1.0.0	100.1.0.13	255.255.255.252	UG	20	0	0
eth3						
100.1.0.4	100.1.0.17	255.255.255.252	UG	20	0	0
eth2						
100.1.0.8	100.1.0.17	255.255.255.252	UG	20	0	0
eth2						
100.1.0.12	0.0.0.0	255.255.255.252	U	0	0	0
eth3						
100.1.0.16	0.0.0.0	255.255.255.252	U	0	0	0
eth2						
100.1.1.0	100.1.0.13	255.255.255.0	UG	20	0	0
eth3						
100.1.2.0	100.1.0.17	255.255.255.0	UG	20	0	0

```
kathara - [ lab: rip_frr ]
```

configuring r4

- step 2: propagating the default route into rip

```
▼ r4
root@r4:/# vtysh

Hello, this is FRRouting (version 8.0.1).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

r4-frr# configure terminal
r4-frr(config)# router rip
r4-frr(config-router)# route 0.0.0.0/0
r4-frr(config-router)# quit
r4-frr(config)# quit
r4-frr# disable
r4-frr> exit
root@r4:/#
```

begin configuration


configure the rip protocol

statically configure the default route

end of rip configuration

end configuration

abandon privileges



the default route


- after a while, the default route has been injected (via rip) into the network

▼ r1

```
root@r1:/etc/frr# route
```

Kernel IP routing table

Destination	Gateway	Genmask	Flags	Metric	Ref	Use
Iface						
default	100.1.0.14	0.0.0.0	UG	20	0	0 eth0
100.1.0.0	0.0.0.0	255.255.255.252	U	0	0	0 eth2
100.1.0.4	100.1.0.2	255.255.255.252	UG	20	0	0
eth2						
100.1.0.8	0.0.0.0	255.255.255.252	U	0	0	0 eth1
100.1.0.12	0.0.0.0	255.255.255.252	U	0	0	0
eth0						
100.1.0.16	100.1.0.10	255.255.255.252	UG	20	0	0
eth1						
100.1.1.0	0.0.0.0	255.255.255.0	U	0	0	0 eth3
100.1.2.0	100.1.0.2	255.255.255.0	UG	20	0	0
eth2						
100.1.3.0	100.1.0.10	255.255.255.0	UG	20	0	0



checking connectivity


▼ r1

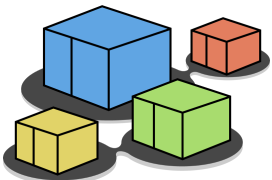
any (even non-existing) destination

```
r1:~# ping 193.204.161.1
PING 193.204.161.1 (193.204.161.1) 56(84) bytes of data.
From 100.2.0.2 icmp_seq=1 Destination Net Unreachable
From 100.2.0.2 icmp_seq=2 Destination Net Unreachable

--- 193.204.161.1 ping statistics ---
2 packets transmitted, 0 received, +2 errors, 100% packet
loss, time 999ms
```

r1:~# █






checking connectivity

- r5 is actually receiving echo request packets

```
r5:~# tcpdump -i eth0 -n -s 1518
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 1518 bytes
11:38:43.822503 arp who-has 100.2.0.2 tell 100.2.0.1
11:38:43.824221 arp reply 100.2.0.2 is-at fe:fd:64:02:00:02
11:38:43.825890 IP 100.1.0.13 > 193.204.161.1: icmp 64: echo request seq 1
11:38:43.827139 IP 100.2.0.2 > 100.1.0.13: icmp 92: net 193.204.161.1
unreachable
11:38:44.841566 IP 100.1.0.13 > 193.204.161.1: icmp 64: echo request seq 2
11:38:44.841651 IP 100.2.0.2 > 100.1.0.13: icmp 92: net 193.204.161.1
unreachable

6 packets captured
6 packets received by filter
0 packets dropped by kernel
r5:~#
```





shutting down an interface

▼ r1



```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
 1 100.1.0.10 (100.1.0.10) 24 ms 1 ms 1 ms
r1:~# ifconfig eth1 down
```



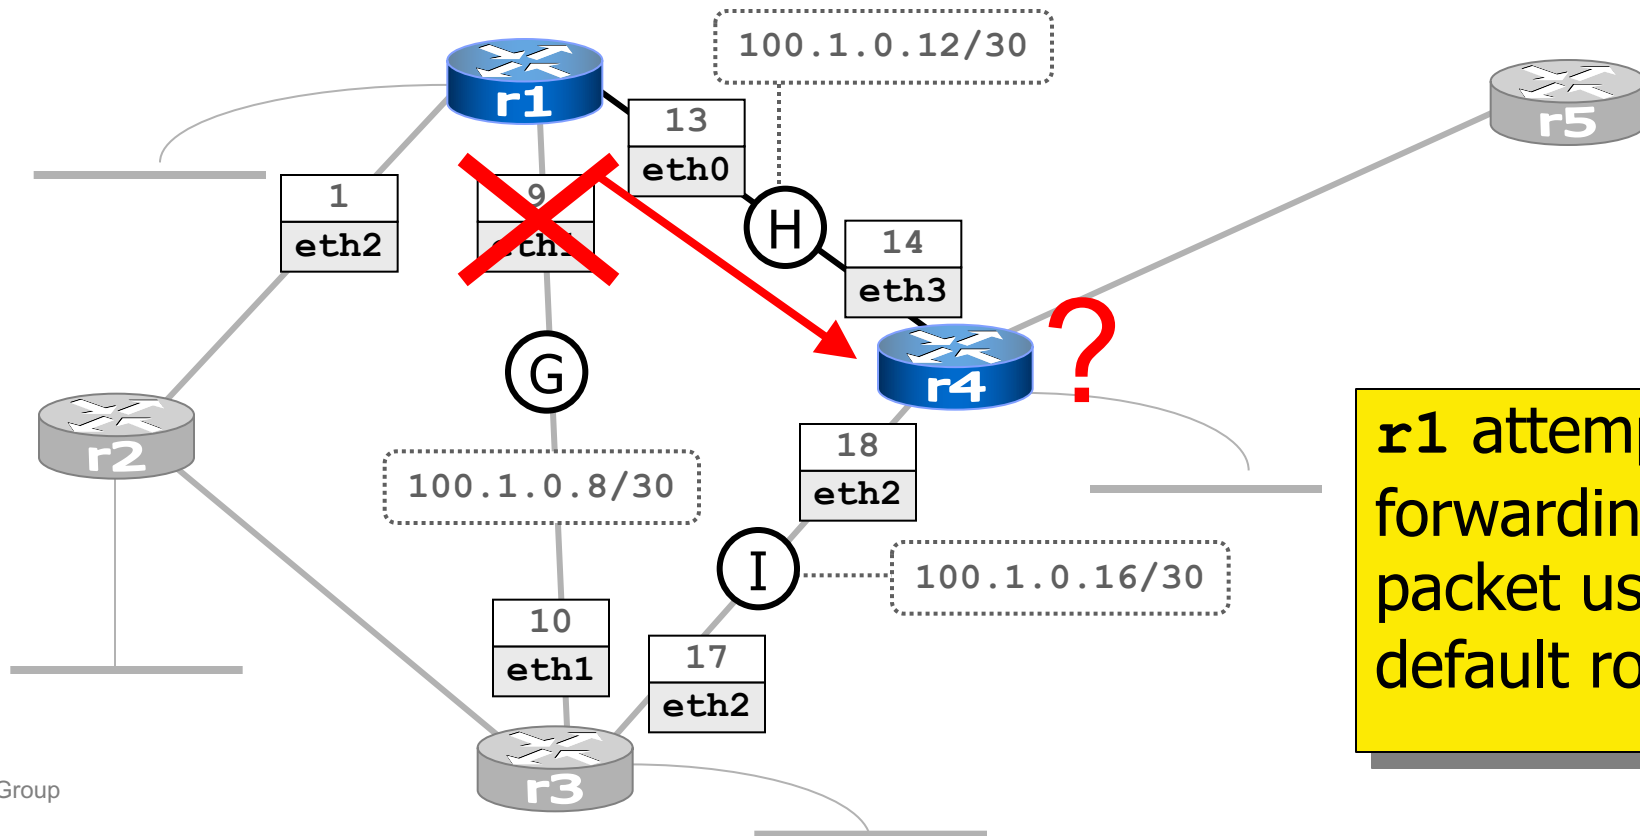


shutting down an interface


▼ r1



```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
 1  100.1.0.14 (100.1.0.14)  1 ms  1 ms  1 ms
 2  * * *
 3  * * * ■
```



r1 attempts forwarding the packet using the default route




shutting down an interface


▼ r1



```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
 1  100.1.0.14 (100.1.0.14)  1 ms  1 ms  1 ms
 2  100.1.0.10 (100.1.0.10)  5 ms  2 ms  1 ms
r1:~# █
```



after a while,
routing
information has
been updated



shutting down an interface

- r1's routing table has been updated

r1

```
r1:~# route
```

Kernel IP routing table

Destination	Gateway	Genmask	Flags	Metric	Ref	Use
100.1.0.16	100.1.0.14	255.255.255.252	UG	2	0	0
eth0						
100.1.0.0	*	255.255.255.252	U	0	0	0 eth2
100.2.0.0	100.1.0.14	255.255.255.252	UG	2	0	0
eth0						
100.1.0.4	100.1.0.2	255.255.255.252	UG	2	0	0
eth2						
100.1.0.8	100.1.0.14	255.255.255.252	UG	3	0	0
eth0						
100.1.0.12	*	255.255.255.252	U	0	0	0 eth0
100.1.4.0	100.1.0.14	255.255.255.0	UG	2	0	0
eth0						
100.1.2.0	100.1.0.2	255.255.255.0	UG	2	0	0
eth2						
100.1.3.0	100.1.0.14	255.255.255.0	UG	3	0	0

kathara - [lab: rip_frr]

last update: Oct 2023